

Chloe Minkyung Kim

Singapore | (+65) 8923 7528 | chloeminkyung@gmail.com | [chloeminkyung.github.io](https://github.com/chloeminkyung)
Core Skills: Java · Python · SQL · Snowflake · ETL & Data Pipelines · Real-time & Batch Architecture

Education

Hong Kong University of Science and Technology

Sep 2016 – Jan 2020

Bachelor of Engineering in Computer Science

Experiences

Goldman Sachs, Singapore

Apr 2024 – Present

Associate

- Leading the design and implementation of the **Lakehouse** platform — a global unified architecture combining the flexibility of a data lake with the performance, structure and governance features of a data warehouse.
- Built core ETL and data transformation and entitlements frameworks mainly in Java, SQL and **Pure**, working with Snowflake, **Legend** and **Alloy** to integrate global business domains.
- Introduced end-to-end data latency improvements: transitioned from multi-day latency to minutes for globally distributed datasets, enabling near-real-time analytics and decision-making.
- Elevated data governance, metadata management and lineage across the platform — helping to establish the Lakehouse as a single source of truth for upstream analytics and downstream applications.
- Positioned the platform for general availability (GA) by end-of-year, coordinating architecture, compliance, operational readiness and cross-regional rollout.

Dozer, Singapore

Dec 2022 – Present

Founding Software Engineer

- Architected and implemented a high-throughput real-time data access platform in Rust, delivering 4.92 ms average latency across almost 10,000 concurrent requests at 50 concurrencies.
- Built streaming connectors (Postgres, Snowflake, Ethereum, Kafka, AWS S3, object stores) and a native SQL engine using RocksDB, enabling joins over 100 million records in 10 minutes.
- Designed and deployed Stripe webhook integration with Rust Diesel, EKS/Minikube deployments, and snapshot/checkpoint mechanism for object stores like S3 and DeltaLake.
- Designed and implemented Stripe webhook with Rust Diesel integration for Dozer cloud deployment service using EKS and Minikube and developed snapshotting checkpointing functionality for object stores such as AWS S3 and Deltalake.
- Presented product demos at Big Data & AI World Singapore, Big Data Conference Europe, and APIdays Paris.

Goldman Sachs, Singapore

Jul 2020 – Dec 2022

Analyst

- Developed and open-sourced a **milestoning component library** using SQL, H2, and Snowflake to version temporal data in large-scale systems.
- Contributed to the Legend open-source data modelling ecosystem via the **Studio tool** and persistence service, improving developer productivity and testing capabilities.
- Improved stability, scalability and latency for the Alloy Streaming ETL platform: leveraged Java, Prometheus, Grafana and Kubernetes to build data-reliability features.
- Operated and optimized Kafka-based distributed ETL pipelines embedded within Kubernetes for high-throughput ingestion and real-time data flows.

JP Morgan Chase & Co., Hong Kong

Jun 2019 – Aug 2019

Summer Analyst

- Built a personalized recommendation system for the Private Banking web platform using Python, Java, Angular and Flask on Gaia (JPMC's cloud), with collaborative filtering and user-clustering algorithms to enhance user experience.

Imago.AI, Hong Kong

May 2018 – Aug 2018

Artificial Intelligence Researcher

- Developed a neural Question and Answer model based on Bi-Directional Attention Flow + self-attention for SQuAD & MS MARCO datasets.
- Created a Chinese Named Entity Recognition model with BiLSTM + CRF targeting Hong Kong address datasets, enhancing extraction accuracy.

Projects & Research Experience

News Data Processing for Fake News Detection, Final Year Project

May 2018 – May 2019

- Designed a multi-modal neural network for fake-news detection, combining text, image, and query-based modules to improve classification accuracy and model diversity.
- Implemented a fact-checking pipeline using BiDAF (Bi-Directional Attention Flow) and AllenNLP Decomposable Attention models, introducing a claim-and-evidence-driven verification approach.
- Deployed trained models behind an exposable REST API / web service, enabling scalable integration with external applications for real-time news verification and content analysis.

BundlePort, Startup Project

Jan 2018 – May 2019

CEO / CTO

- Founded and led a startup developing a full-stack web platform for retail logistics management, built with Node.js and MySQL, and deployed on AWS EC2, S3, Elastic Beanstalk, and CloudFront.
- Directed overall product vision, engineering roadmap, and sales operations, managing cross-functional coordination between development and retail teams.
- Focused on UI/UX optimization and process automation, improving transaction efficiency and streamlining logistics workflows for partner merchants.

Skills

Technical Skills	Java, Python, Java, Python, SQL, Snowflake, Rust, PyTorch, Tensorflow, C++, Prometheus
Language Skills	English (Fluent), Korean (Native), Mandarin (Intermediate)

Honors and Awards

- JPMorgan Chase & Co. Internal Innovation Week Hackathon 2019 Winner (JPMC Trends), 2019
- SmarTone Hackathon 2018 "Smart Properties", Interactive Property Chatbot - Top 3 Project, 2018
- Targeted Scholarship under HKSAR Government Scholar Fund, Full Scholarship, 2016 - 2020
- Admission Scholarship from HKUST, Full Scholarship, 2016 - 2020
- Mensa Membership from Malaysia Mensa (Mensa International), High IQ society, 2015